

 Approved Startup #1924 • Winner of PAA AI in Business & Industry Award

The Sovereign, Self-Hosted Autonomous AI Workforce for Enterprises & Government

Fathom delivers universal digital employees deployed inside your private infrastructure. Powered by a high-performance compiled Rust server engine and an intuitive desktop chat application, executing end-to-end workflows with full auditability, computer use, and team coordination.

SOVEREIGN DEPLOYMENT

100% On-Premise

Zero vendor data leaks. Runs air-gapped on bare-metal, private cloud, or sovereign state data centers.

RUST CORE SPEED

~0.75 ms Dispatch

15.4 MB memory baseline, sub-millisecond tool routing, and Tokio asynchronous sub-agent swarms.

FATHOM BOT SURFACE

Local-First Chat App

Every contact is an autonomous agent. Multi-agent group chats, routine cron schedules, and computer control.

ENTERPRISE TRUST

Human-in-the-Loop

Signed verification receipts, strict permission cards, and RBAC governance built for regulated sectors.

The Enterprise AI Conundrum: Cloud Leaks vs. Toy Scripts

Companies and governments want autonomous digital workers, but current options fail three critical enterprise requirements.



Data Sovereignty & Security Leaks

Public AI SaaS (OpenAI, Anthropic Cloud, Copilot) force enterprises to stream proprietary banking, customer, and state data to foreign servers. Regulators strictly penalize cross-border personal data transfers.

82% of enterprises restrict public cloud AI due to compliance risk.



Bloated Python Stacks & Fragile Scripts

Frameworks like LangChain, CrewAI, and AutoGen are written in interpreted Python. They consume 2.5 GB RAM per agent, suffer from GIL bottlenecks, ~400ms latency, and crash during multi-hour background tasks.

Unstable in production; incapable of running 10+ sub-agents reliably.



Uncontrollable Black Boxes Without Approvals

Existing autonomous tools lack granular human governance. They execute unchecked terminal commands, mutate files, and interact with external systems without cryptographic audit trails or interactive approval cards.

Executives refuse full autonomy without deterministic safety rails.

The Solution: Full-Stack Sovereign Autonomous Workforce

Fathom combines an enterprise on-premise server engine with an intuitive desktop chat workspace where every colleague is an autonomous specialist.

CORE ENGINE

1. Compiled Rust Runtime

Built with Tokio asynchronous swarms, sub-millisecond dispatch (~0.75 ms), 15.4 MB memory baseline, SQLite/PostgreSQL persistence, and long-term semantic memory.

- Zero Python dependencies in runtime core
- 13 modular workspace crates
- 51 always-registered built-in tools

OPERATOR SURFACE

2. Fathom Bot Chat App

Local-first desktop chat client (React 19 + Electron). Talk to agents in 1-on-1s, add multiple bots to shared discussion channels, and schedule autonomous routines via visual cron.

- Every sidebar contact is a specialized agent
- Multi-bot group discussions & peer reviews
- Calendar automations & morning briefings

EXECUTION & SAFETY

3. Governed Computer Use

Agents control browser tabs and VM desktops with visual feedback. High-risk actions require explicit human click-to-approve cards before execution.


- Interactive human-in-the-loop approval cards
- Signed cryptographic verification receipts
- Sandboxed Linux/macOS VM environments

Enterprise Ready Today: Connects to amoCRM, Bitrix24, 1C, SAP, private databases, and sovereign local LLMs without external dependencies.

100% PRIVATE & SOVEREIGN

Universal Autonomous Digital Coworkers for Any Enterprise Mission

Operators configure custom digital employees with distinct identities, system directives, tool access boundaries, specialized long-term memory, and local or private LLM backends.



24/7 Autonomous Operations: Scheduled Background Routines & Triggers

Digital coworkers independently execute multi-hour scheduled workflows: daily market intelligence summaries, automated customer audits, and continuous system monitoring without human supervision.

F

FATHOM

SYS

+

Q

Search

Bots

▼

>

Tokyo Weather Assistant

Turnip

2:53 AM

System Resources

Q

fathom

Turnip

CoH

1

140k

Turnip wants to other

uptime && cat /proc/loadavg

✓ Allowed

• System Resources

CPU

10 cores (Apple Silicon)

Load Average

1.28 1m / 1.10 (5m)

Memory Used

1.9 / 7.8 GiB

24%

Available Memory

5.9 GiB

Jump to latest

• CPU: 10 virtual cores (Apple Silicon architecture, base 2.0 GHz)

• Memory: 7.8 GiB total — 1.9 GiB used (24%), 5.9 GiB available

• Swap: 8.8 GiB total (20 MiB used)

• Load Average: 1.28 (1m), 1.10 (5m), 0.90 (15m)

This run

3 steps

✓ lscpu

lscpu || uname -a

✓ free

free -h

✓ uptime

uptime && cat /proc/loadavg

Save as skill

Adds the run to your message; add any notes, then send.

Tools

↑

SYS.ROUTER

18ms

You

⌘1 Статус серверов

⌘2 Безопасность

⌘3 Docker

⌘4 Аналитика

Ask


Message Turnip

Thread workspace

Tools

Autonomous Tool Orchestration with Transparent Chain-of-Thought

Every agent autonomously breaks high-level goals into multi-step execution plans, dynamically selects tools (shell, web, files, CRM), and provides transparent real-time reasoning streams.



Deterministic Multi-Step Execution with Verifiable Proof of Work

Zero black-box hallucinations. Complex tasks are decomposed into audited execution stages with signed receipts, explicit environment validation, and step-by-step progress confirmation.

The image shows a dark-themed user interface for an AI application named "FATHOM".

On the left is a sidebar with:
- A top section with "F FATHOM" and "SYS" tabs, a search bar, and a "Bots" dropdown.
- A list of bots, with "Turnip" selected and showing a timestamp of "2:48 AM".
- A "Tools" section at the bottom with a "SYS.ROUTER" entry and a "You" entry.

The main area is a chat window titled "Turnip". It shows a message from the bot with an "INTERNAL THOUGHT STREAM" section. The text in this stream indicates the bot is establishing its identity, analyzing the user's intent, and asking for confirmation to proceed with setup. Below the thought stream, the bot says "Thinking" with a duration of "3m 04s".

On the right side of the chat, there is a "This run" summary panel showing a list of 4 steps:
1. exa "exa key"
2. Turnip "Turnip setup"
3. thread "thread cc88bedc-3b4b-49e4-b4c2-b6..."
4. exa "exa"

At the bottom of the interface, there is a bar with a "Thread workspace" button and a "Tools" button. The user's input "Weather in Tokyo" is visible in the chat area.

Autonomous Web Exploration, OSINT & Structured Enterprise Synthesis

Agents crawl external registries, competitor databases, and industry websites to synthesize live market intelligence into structured, decision-ready executive tables and reports.

F FATHOM SYS

Search

Bots

Tokyo Weather Assistant

Turnip 2:58 AM

Uzbekistan VC & Investment Fun...

Tools

SYS.ROUTER 18ms

You

localhost

Turnip

Uzbekistan VC & Investment Funds

UzVC (National Venture Fund) State Fund Fund of Funds

AloqaVentures Corporate VC Seed / Pre-seed

Sturgeon Capital Regional VC Seed / Series A

MOST Ventures Central Asia Early Stage

IT Park Uzbekistan Ecosystem Ventures / Grants

UzAngels Network Angel Club

Key venture funds and investment organizations active in Uzbekistan

- UzVC: info@uzvc.uz | uzvc.uz — National venture capital fund.
- AloqaVentures: info@aloqaventures.uz | aloqaventures.uz — VC backed by Aloqabank.
- Sturgeon Capital: info@sturgeoncapital.com | sturgeoncapita VC actively backing Uzbek startups.
- MOST Ventures: dealflow@most.com.kz | most.com.kz — Regional fund & syndicate

Tools

Ask Message Turnip

Thread workspace Tools

This run 5 steps

Uzbekistan "Uzbekistan venture capital fu...

Uzbekistan "Uzbekistan venture capital fu...

UzVC "UzVC venture capital"

UzVC "UzVC"


Uzbekistan "Uzbekistan VC"

Save as skill

Adds the run to your message; add any notes, then send.

Enterprise Infrastructure Diagnostics & Resource Governance

Autonomous monitoring of host servers, container health, CPU/RAM utilization, and token throughput — enabling agents to diagnose system bottlenecks and report issues proactively.



Digital coworkers don't just reply with raw text. They deliver interactive dashboards, live weather reports, clickable action cards, and rich mini-apps directly inside the team discussion.

⬅️📅Automations

ScheduleRun logsWebhooks

☰List📅Calendar

⬅️Today➡️

September 7 – 13, 2026

September 2026

M T W T F S S

31 1 2 3 4 5 6

7 8 9 10 11 12 13

14 15 16 17 18 19 20

21 22 23 24 25 26 27

28 29 30 1 2 3 4

5 6 7 8 9 10 11

👤MY BOTS1

🔍Search bots

⋮

Turnip

Tokyo Weather Assistant

GMT+3

MON7

TU8

1 AM

2 AM

3 AM

4 AM

5 AM

6 AM

7 AM

8 AM

9 AM

New calendar event

✕

Add title

RoutineCall

🕒Tuesday, September 8
2:00 AM

👤+Turnip

📄What should the bot do?

Post results to

Create a dedicated results thread

Each run starts with fresh context. Dated results collect here; full execution details stay in Run logs. A dedicated results thread is reused for every run and can be moved into a folder.

More optionsSave

1

All bots

Week

SAT12

SUN13

Почасовой мо...
1:54 AM · Every h

Drag a bot onto any time to schedule it.

The Inflection Point: From Conversational Chat to Agentic AI

Global analyst consensus confirms that enterprise spending is decisively moving away from passive chatbots toward autonomous, task-executing agent swarms.

\$1.3T+

IDC: AGENTIC AI SPENDING BY 2029

IDC projects Agentic AI will exceed 26% of all worldwide enterprise IT spending within 5 years, marking the industry's next generational compute transition.

40%

GARTNER: ENTERPRISE APPS WITH AGENTS

Gartner forecasts 40% of all enterprise applications will embed task-specific autonomous AI agents, jumping from under 5% in early 2025.

\$5.95T

GARTNER: GLOBAL AI SPENDING BY 2030

Total enterprise AI investment surging to nearly \$6T as traditional human knowledge workflows transition to persistent digital workers.

Analyst Takeaway: The "Toy Chatbot" Era Has Ended

Enterprise buyers no longer allocate budget for novelty chat windows that output text advice. They demand sovereign digital agents that execute complete multi-step tasks inside their existing CRM, ERP, and database perimeter.

GARTNER & IDC CONSENSUS

The New Economics of Labor: McKinsey Agentic Advantage

Deploying sovereign autonomous workers shifts enterprise cost structures from linear human payroll to sub-dollar compute execution.

MGI RESEARCH

60%–70% Labor Hour Automation

McKinsey Global Institute calculates that current generative and agentic AI capabilities can automate corporate work activities absorbing 60% to 70% of employee time today.

COST DEFLATION

92% Cost Reduction per Task

A junior analyst task costing \$35/hr in human labor (research, CRM data entry, outreach drafting) executes on Fathom in ~40 seconds for less than \$0.08 in local server compute.

ORGANIZATIONAL VELOCITY

10× Execution Throughput

Digital coworkers operate 24/7 without context switching or fatigue. Back-office queues in banking, tax auditing, and vendor vetting shrink from weeks to minutes.

Enterprise Invariant: Sovereign Security Unlocks Deployment

Enterprises cannot realize these 90%+ productivity gains if using cloud tools that violate state data regulations. Fathom's 100% on-premise Rust engine delivers full agentic labor economics inside corporate firewalls.

MCKINSEY RESEARCH BACKED

Addressable Market: Global \$165B Wave with Sovereign Focus

Massive global tailwinds paired with an immediate, high-margin niche in sovereign state institutions and regulated enterprise clusters.

\$165B

TAM: GLOBAL ENTERPRISE AI AGENT MARKET (BY 2032)

43.8% CAGR driven by enterprise back-office automation, legal, finance, and industrial workflows.

\$12.4B

SAM: SOVEREIGN ON-PREMISE & REGULATED GOV/B2B

Financial institutions, telecom operators, healthcare, and state ministries requiring sovereign data silos.

\$180M

SOM: CENTRAL ASIA, MENA & EMERGING TECH HUBS

Initial 3-year serviceable market across Uzbekistan, Kazakhstan, UAE, and GCC enterprise digitization programs.

Macro Drivers Accelerating Sovereign Adoption:

- 1. Open-Weight Model Revolution:** Open models now match closed proprietary APIs, making private on-premise execution superior in cost, latency, and data control.
- 2. Strict Regulatory Mandates:** Personal Data Protection laws (e.g. Uzbekistan Law #547, EU AI Act) mandate that sensitive citizen and corporate records remain strictly on sovereign soil.

43.8%

ANNUAL INDUSTRY CAGR

The Sovereign Mandate: Compliance as an Enterprise Moat

Global data localization laws are erecting permanent barriers against foreign cloud AI, creating massive demand for self-hosted enterprise agent engines.

CENTRAL ASIA

Uzbekistan Law #547 on Personal Data

Mandates physical database localization within the Republic of Uzbekistan for processing citizen records. Foreign cloud AI tools create direct legal liability for corporate executives.

EUROPEAN UNION

EU AI Act & GDPR Strictures

Bans unmonitored automated biometric, financial, and workplace decision-making. Requires deterministic logging, human-in-the-loop overrides, and verifiable audit trails.

MIDDLE EAST / GCC

UAE & KSA Sovereign Cloud Policies

Government entities and national banks must deploy AI within national sovereign data centers (e.g. Injazat, Moro Hub), prohibiting transmission to offshore LLM APIs.

Fathom's Native Architecture Fits Existing Compliance Frameworks Out of the Box:

- **Zero External Phoning:** Can run completely air-gapped without an active Internet connection using local quantized LLM weights.
- **Signed Verification Receipts:** Cryptographically records every tool call, file write, and bash command into an unalterable local audit journal.

Why Fathom Wins: Compiled Rust Performance & Sovereignty

How Fathom fundamentally outperforms public cloud bots, legacy Python agent frameworks, and enterprise copilots.

OUR PLATFORM  FATHOM Sovereign AI Workforce Deployment: 100% Self-Hosted & Air-Gapped Engine: Compiled Rust 1.97 (~0.75 ms dispatch) Memory: 15.4 MB baseline RAM footprint Interface: Fathom Bot (React 19 + Electron) Collaboration: Multi-bot channels & group critiques Governance: Signed receipts & approval cards	CONSUMER CLOUD BOTS Grok Bot / Google Spark Proprietary Web Chat Deployment: Public Cloud Only (Vendor Lock-in) Engine: Proprietary Cloud Cluster Memory: Cloud-hosted (No on-prem option) Interface: Web / Mobile Single Chat Collaboration: 1-on-1 chat only (No team channels) Governance: Opaque cloud telemetry logs	OPEN SOURCE PYTHON CrewAI / LangChain Developer Scripts Deployment: Self-hosted Python Scripts Engine: Interpreted Python (~400ms lag) Memory: 1.8 GB – 2.5 GB RAM per agent Interface: Terminal CLI / Jupyter Notebooks Collaboration: Rigid linear sequential chains Governance: Raw stdout logs without receipts	OFFICE COPILOTS M365 Copilot Sidebar Document Assistant Deployment: Multi-tenant SaaS Cloud Engine: Closed Cloud Microservices Memory: N/A (Closed SaaS) Interface: Office Sidebar Plugin Collaboration: Single-turn document autocomplete Governance: Cloud enterprise admin portal
--	---	---	--

Transparent B2B Monetization with Unlimited AI Tokens

Simple, predictable pricing: personal desktop operators and custom enterprise deployments with zero per-token anxiety.

INDIVIDUAL OPERATORS

Personal Operator

\$79 / month

 **Unlimited AI Tokens Included**

- ✓ Full Fathom Bot desktop application (macOS, Windows, Linux)
- ✓ Unlimited custom agent personas & behavioral tuning
- ✓ Autonomous routines & background calendar scheduler
- ✓ Deep web research agent & OSINT tools
- ✓ Standard CRM connectors (amoCRM, Bitrix24)
- ✓ Community support & continuous updates

ENTERPRISE & GOV SOVEREIGN

Enterprise Sovereign

Custom (\$200 – \$500 / mo / node)

 **Unlimited AI Tokens • 100% On-Premise Air-Gapped**

- ✓ 100% On-Premise deployment inside private infrastructure
- ✓ Unlimited internal operator seats across the organization
- ✓ Multi-bot shared team channels with peer review delegation
- ✓ Governed computer use with interactive permission cards
- ✓ Custom ERP & intranet integration (1C, SAP, private SQL)
- ✓ Hardware acceleration support (CUDA / Apple Metal)
- ✓ Enterprise SLA, dedicated security patching & onboarding

84%

SOFTWARE GROSS MARGIN

Zero

TOKEN ANXIETY FOR CLIENTS

11.8x

PROJECTED LTV / CAC RATIO

From Breakthrough Rust Engine to Enterprise Fleet Scale

Proven technology, approved government startup accreditation, and a clear path to regional leadership.

Q2-Q3 2026 • COMPLETED

Core Engine & MVP

- 13 Rust workspace crates built
- 51 built-in tools + computer control
- Fathom Bot React/Electron client
- **Winner of PAA AI Award (#1924)**
- 42-page technical whitepaper published

Q4 2026 • CURRENT

Commercial Pilots

- 5 enterprise pilots deployed
- amoCRM, Bitrix24 & 1C native sync
- Multi-bot channel collaboration
- Sovereign on-premise bundle
- 500+ active waitlist operators

Q1-Q2 2027 • EXPANSION

Regional Growth

- Scale to 40+ paid B2B clients
- Government ministry pilot in Uzbekistan
- Kazakhstan & UAE partner channels
- Expanding recurring revenue fleet
- SOC 2 Type II certification

Q3-Q4 2027 • SCALE

Sovereign Fleet

- Expansion into GCC banking sector
- Fathom Marketplace for custom bots
- Fine-tuned local domain SLMs
- Scaled enterprise node fleet
- Institutional funding expansion

 **Accredited by StartupBase & IT Park Uzbekistan**
Startup ID: #1924 • Category: AI in Business & Industry (Sanoat va tadbirkorlikdagi SI) • OneID Verified Founders

STAGE 3: APPROVED BY COMMITTEE

Deep Technical Expertise in High-Performance Systems & Growth

Founders combining low-level systems engineering, enterprise product execution, and regional B2B growth.



German Yakushev

FOUNDER & CEO • LEAD ARCHITECT

Systems programmer and technology founder. Designed and built Fathom's 13-crate Rust engine, Tokio async multi-agent pipeline, and Fathom Bot messaging architecture. Deep experience in autonomous agent systems and enterprise software architecture.

📍 Tashkent, Uzbekistan • PINFL: 31210968380014



Mirtemir Anorboyev

CO-FOUNDER & PRODUCT MANAGER

Product strategist directing operator user experience, enterprise workflow discovery, and automated templates. Oversees product roadmap alignment between technical architecture and customer requirements across banking, CRM, and corporate back-office workflows.

📍 Tashkent, Uzbekistan • PINFL: 51607005470010



Elbek Yuldashev

CO-FOUNDER & HEAD OF GROWTH

Enterprise sales and growth lead managing commercial pilots, government agency relations, and channel partnerships across Central Asia and the CIS. Spearheads enterprise outbound acquisition, customer success onboarding, and corporate pilot conversions.

📍 Tashkent, Uzbekistan • PINFL: 51204066180078

Scaling the Sovereign AI Workforce Across Central Asia & MENA

Partnering with visionary investors to accelerate enterprise deployments, state certifications, and commercial fleet scaling.

INVESTMENT PROPOSITION

Seed Capital for Strategic Scaling

We are opening an early-stage investment round for strategic partners to accelerate commercial enterprise pilots, establish sovereign government installations, and expand engineering capacity over an 18-month growth runway.

- ✓ **Instrument:** Standard SAFE / Convertible Promissory Note
- ✓ **Runway:** 18 Months of aggressive product & sales acceleration
- ✓ **Key Milestone:** 30+ Enterprise Sovereign Nodes deployed

STRATEGIC MILESTONE TARGET

Expansion across Banking, Government & Telecommunications

Use of Capital Allocation

Systems Engineering & Fine-Tuning 45%

Senior Rust engineers, local model quantization (vLLM/Ollama), and VM computer control virtualization.

Enterprise Sales & Government Relations 30%

B2B enterprise account executives, pilot implementation engineers, and regional summit sponsorships.

Security, Certifications & Legal 15%

State personal data compliance, SOC 2 preparation, penetration testing, and corporate IP registration.

Infrastructure & Testing GPU Nodes 10%

Private GPU cluster for automated agent regression benchmarks and local model evaluation.

Deploy Autonomous AI Workforce Without Compromising Data Sovereignty

We are ready to demonstrate a live deployment on your private infrastructure in under 15 minutes. Join us in establishing the sovereign AI backbone for businesses and governments.



German Yakushev (CEO)

+998 90 061 80 26

yakushev00@gmail.com

Telegram: @yakushev_hk



Live Product Resources

Web Platform: fathom.uz

Fathom Bot App: fathom.uz/bot

Systems Whitepaper: fathom.uz/whitepaper

GitHub Repository: github.com/yakushevkh/Fathom



Legal & Accreditation

StartupBase ID: #1924 (Approved)

Track: AI in Business & Industry

Entity: Technology X / Fathom

Location: Tashkent, Uzbekistan

Building the future of sovereign enterprise intelligence together.