

Unlimited Neural Compute Engine

Harnessing Frontier Foundation Models for High-Throughput Cost Arbitrage

PAGE 07 / 42

The Foundation Model Cost-Performance Revolution

Offering unlimited neural compute is economically viable because Fathom intelligently routes tasks to next-generation frontier foundation models that deliver elite reasoning at a fraction of legacy pricing.

KIMI K3 (MOONSHOT AI)

Strength: Ultra-Long Context

Processes massive document corpora, multi-page regulatory filings, and deep recursive research trees with flawless long-range recall across 200k+ tokens.

QWEN 3.8 MAX (ALIBABA)

Strength: Tool Calling & Code

Exceptional precision in structured function calling, multilingual web parsing, and Python/Node.js REPL script generation with microsecond JSON serialization.

GLM 5.3 (ZHIPU AI)

Strength: Fast Decomposition

High-speed reasoning engine optimized for coordinator agents breaking down complex user briefs into parallel worker subtasks with sub-200ms TTFT.

Fathom's Economic Arbitrage Flywheel

1. Compiled Rust Core

~15MB RAM & 0.75ms dispatch



2. Efficient Chinese LLMs

Kimi k3 / Qwen / GLM



3. Sub-\$25 Monthly Cost

Total compute per worker



4. 90%+ Gross Margin

On flat seat subscription

DYNAMIC ROLE-BASED MODEL ROUTING

The coordinator dynamically assigns the ideal model per role: GLM 5.3 for high-speed planning, Kimi k3 for deep document synthesis, and Qwen 3.8 Max for browser automation and code execution. This maximizes accuracy while keeping internal token expenditures at micro-cents per task.

The Economic Reality: Token costs on these frontier engines range between \$0.10 and \$0.40 per Million tokens—over 15x cheaper than legacy Western API models—delivering 90%+ gross margins on flat seat subscriptions.