

Why Rust? The Architecture of Performance

Zero-Cost Abstractions, Memory Safety & The Elimination of Python Bottlenecks

The Engineering Foundation: Why We Chose Rust

Traditional agent frameworks built on Python suffer from high latency, heavy memory consumption, fragile type errors, and concurrency bottlenecks caused by the Global Interpreter Lock (GIL). **Fathom is built natively in Rust** to provide enterprise-grade reliability and speed.

THE PYTHON FRAMEWORK TRAP

- **Massive Memory Footprint:** Idle Python runtimes consume 400MB–1.5GB RAM per agent process.
- **GIL Concurrency Deadlocks:** Asynchronous I/O is serialized, choking multi-agent swarms.
- **Slow Startup Latency:** 2.5 to 8.0 seconds just to load interpreter and dependency trees.

THE FATHOM RUST ADVANTAGE

- **Microscopic RAM Footprint:** Lean 15–35 MB baseline RAM allows 100+ agents per server.
- **True Multi-Core Parallelism:** Tokio async tasks utilize all CPU cores without locks.
- **Instant Binary Startup:** Starts in under 5 milliseconds from cold execution.

Architectural Efficiency Comparison

ENGINEERING METRIC	PYTHON AGENT FRAMEWORKS	FATHOM RUST RUNTIME
Tool Dispatch Latency	25.0 – 150.0 ms	0.75 ms (Microsecond Level)
Memory Usage (100 Agents)	40 – 120 GB RAM (Requires Cluster)	1.5 – 3.5 GB RAM (Single Modest VM)
HTML Parsing Speed	15,000 rows/sec (BeautifulSoup)	350,000+ rows/sec (Scraper Rust)
Binary Packaging	Fragile venv + wheels	Single Static Binary (Zero Dependencies)

Speed Equals Intelligence: Microsecond tool execution means agents spend 99.9% of their time waiting for model tokens, not framework overhead.